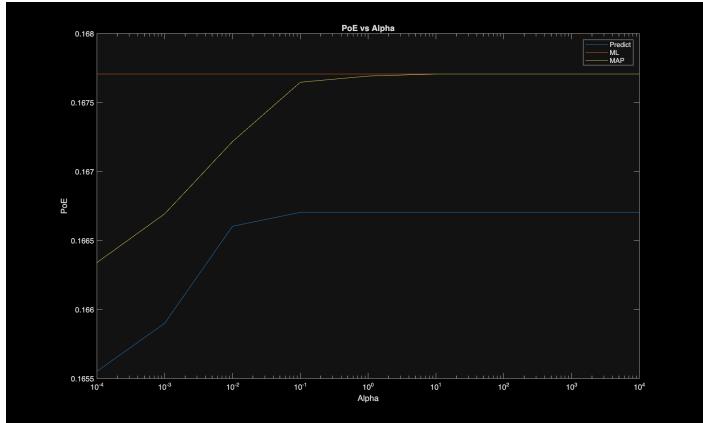


ECE 271A Homework 3

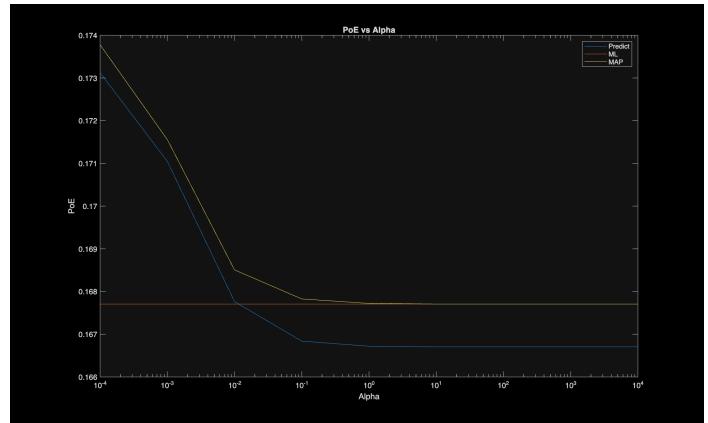
Raghusrinivasan Venkatesan (A69044562)

Figures

Dataset 1

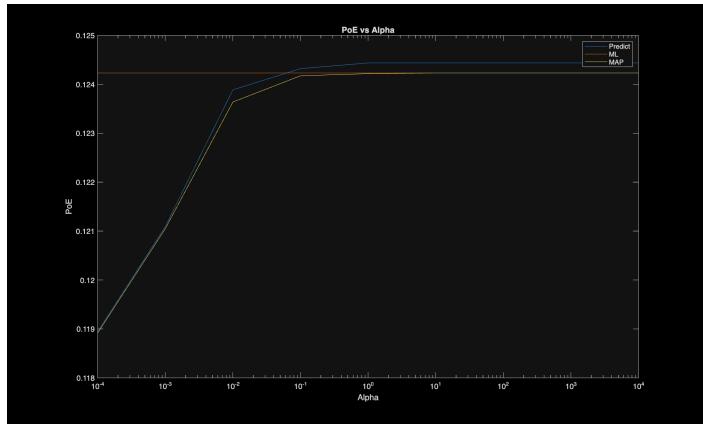


Strategy 1

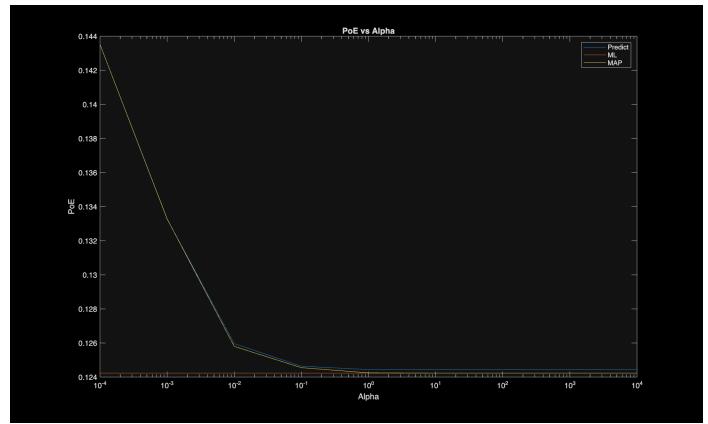


Strategy 2

Dataset 2

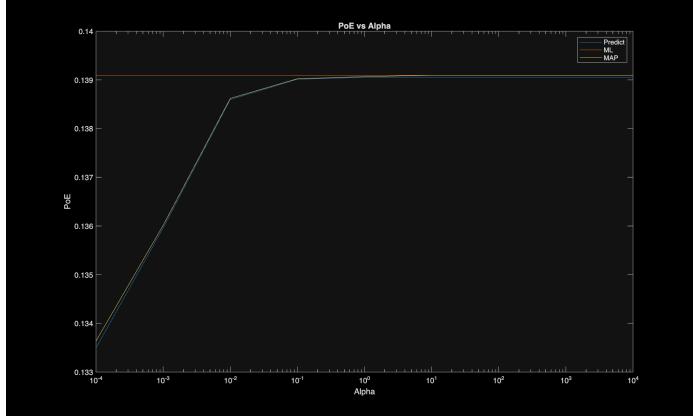


Strategy 1

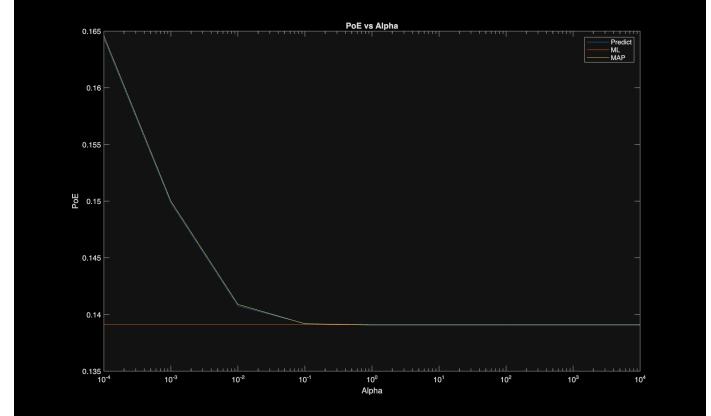


Strategy 2

Dataset 3

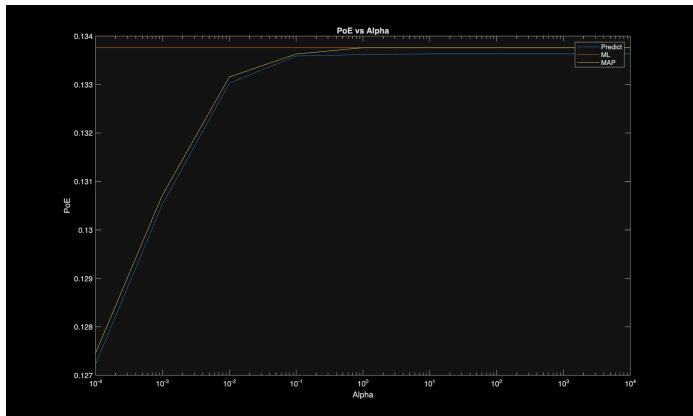


Strategy 1

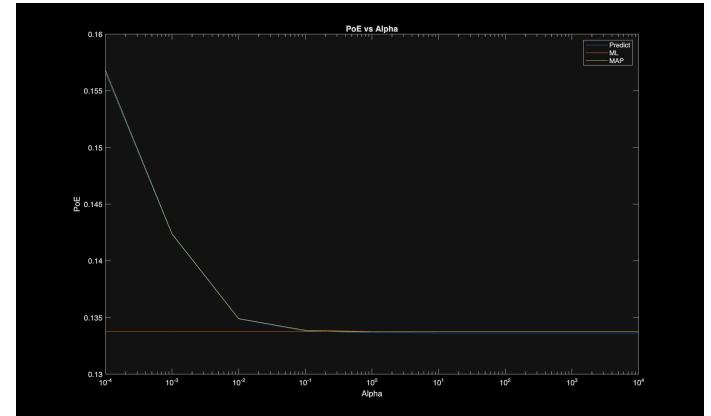


Strategy 2

Dataset 4



Strategy 1



Strategy 2

Part 1: The relative behavior of these three curves

Observation: The probability of error (PoE) exhibits distinct patterns across the three classification methods as the regularization parameter alpha increases. At low alpha values, the three curves typically show significant separation, with Bayes and MAP demonstrating different error rates compared to ML. As alpha increases, the curves tend to converge and stabilize at higher alpha values, with relative performance differences diminishing.

Explanation:

Let $\Sigma_0 = \alpha W_0$ (prior covariance scaling with alpha).

Rewrite posterior mean using weights: $\mu_n = w_{\text{data}} \times \hat{\mu} + w_{\text{prior}} \times \mu_0$

where: $w_{\text{data}} = \Sigma_0(\Sigma_0 + (1/n)\Sigma_n)^{-1}$ $w_{\text{prior}} = (1/n)\Sigma_n(\Sigma_0 + (1/n)\Sigma_n)^{-1}$

Note: $w_{\text{data}} + w_{\text{prior}} = I$ (weights sum to identity)

Case 1: Small α (Small Prior Covariance)

When $\alpha \rightarrow 0$, $\Sigma_0 \rightarrow 0$:

$$T = (\Sigma_0 + (1/n)\Sigma_n)^{-1} \approx ((1/n)\Sigma_n)^{-1} = n\Sigma_n^{-1}$$

$$w_{\text{data}} \approx \Sigma_0 \cdot n\Sigma_n^{-1} \rightarrow 0 \quad w_{\text{prior}} \approx (1/n)\Sigma_n \cdot n\Sigma_n^{-1} = I$$

Result: At small α , posterior mean dominated by prior

$\mu_n \approx \mu_0$ (prior mean dominates)

Comparison of Methods at Small α :

- ML Method: $\mu_{\text{ML}} = \hat{\mu}$ (only uses sample data)
- MAP Method: $\mu_{\text{MAP}} = \mu_n \approx \mu_0$ with $\Sigma_{\text{MAP}} = \Sigma_n$
- Bayes Method: $\mu_{\text{Bayes}} = \mu_n$ with $\Sigma_{\text{Bayes}} = \Sigma_n + \Sigma^{\text{post}}$ (larger covariance)

Conclusion: $|\mu_{\text{ML}} - \mu_{\text{MAP}}|$ is large \Rightarrow Large PoE difference between methods

Case 2: Large α (Large Prior Covariance)

When $\alpha \rightarrow \infty$, $\Sigma_0 \rightarrow \infty$:

$$T = (\Sigma_0 + (1/n)\Sigma_n)^{-1} \approx \Sigma_0^{-1} = (1/\alpha)W_0^{-1}$$

$$w_{\text{data}} = \Sigma_0 T \approx \alpha W_0 \cdot (1/\alpha)W_0^{-1} = I \quad w_{\text{prior}} = (1/n)\Sigma_n T \approx 0$$

Result: At large α , all methods use similar mean estimates

$\mu_n \approx \hat{\mu}$ (data-driven estimate dominates)

Also: $\Sigma^{\text{post}} = \Sigma_0 T (1/n) \Sigma_n \approx (1/n) \Sigma_n \rightarrow 0$ as $\alpha \rightarrow \infty$

Comparison of Methods at Large α :

- ML Method: $\mu_{\text{ML}} = \hat{\mu}$
- MAP Method: $\mu_{\text{MAP}} \approx \hat{\mu}$ with $\Sigma_{\text{MAP}} = \Sigma_n$
- Bayes Method: $\mu_{\text{Bayes}} \approx \hat{\mu}$ with $\Sigma_{\text{Bayes}} \approx \Sigma_n$

Conclusion: All three methods converge to same mean and similar covariance \Rightarrow All methods have similar PoE (convergence observed in plots)

Part 2: How Behavior Changes Across Different Datasets

Observation: The convergence rate and relative performance of the three methods vary significantly across datasets. Datasets with smaller sample sizes (D1, D2) show more pronounced method differences and steeper PoE curves at low alpha values. Datasets with larger sample sizes (D3, D4) demonstrate faster convergence and more modest PoE variations across the alpha spectrum.

Explanation:

The relative influence of prior versus data is:

$$\text{Influence Ratio} = (\text{Prior strength}) / (\text{Data strength}) = \Sigma_0^{-1} / (n\Sigma_n^{-1})$$

$$\begin{aligned} \text{In the posterior update, rewrite as: } \mu_n &= (\Sigma_0^{-1} + n\Sigma_n^{-1})^{-1}(\Sigma_0^{-1}\mu_0 + n\Sigma_n^{-1}\hat{\mu}) \\ &= [\Sigma_0^{-1}/(\Sigma_0^{-1} + n\Sigma_n^{-1})] \times \mu_0 + [n\Sigma_n^{-1}/(\Sigma_0^{-1} + n\Sigma_n^{-1})] \times \hat{\mu} \end{aligned}$$

Define weights:

$$w_{\text{prior}} = \Sigma_0^{-1}/(\Sigma_0^{-1} + n\Sigma_n^{-1}) = 1 / (1 + n\Sigma_n^{-1}\Sigma_0)$$

$$w_{\text{data}} = n\Sigma_n^{-1}/(\Sigma_0^{-1} + n\Sigma_n^{-1}) = n\Sigma_n^{-1}\Sigma_0 / (1 + n\Sigma_n^{-1}\Sigma_0)$$

Case 1: Small Sample Size ($n \rightarrow$ small)

Example from code: $n_{\text{FG}} = 100$, $n_{\text{BG}} = 100$

$$w_{\text{prior}} = 1 / (1 + 100 \times \Sigma_n^{-1}\Sigma_0)$$

Since $100 \times \Sigma_n^{-1}\Sigma_0$ is relatively small: $w_{\text{prior}} \approx$ significant (non-negligible) $w_{\text{data}} \approx$ smaller

Result: Prior information remains influential even at moderate α values

Conclusion: Methods differ significantly across entire α range \Rightarrow Steep PoE curves with large variations (observed in D1, D2)

Case 2: Large Sample Size ($n \rightarrow$ large)

Example: $n_{\text{FG}} = 1000$, $n_{\text{BG}} = 1000$

$$w_{\text{prior}} = 1 / (1 + 1000 \times \Sigma_n^{-1}\Sigma_0)$$

Since $1000 \times \Sigma_n^{-1}\Sigma_0$ is much larger than Case 1: $w_{\text{prior}} \approx 0$ (nearly negligible) $w_{\text{data}} \approx 1$ (dominates)

Result: Data dominates quickly, prior influence vanishes rapidly

Conclusion: Methods converge quickly to ML behavior at small $\alpha \rightarrow$ Flatter PoE curves with fast convergence (observed in D3, D4)

Quantitative Comparison:

$$\text{Convergence factor} = [w_{\text{prior}}(n=100)] / [w_{\text{prior}}(n=1000)] = [1 + 1000 \times \Sigma_n^{-1} \Sigma_0] / [1 + 100 \times \Sigma_n^{-1} \Sigma_0] > 1$$

Convergence Rate $\propto O(1/n)$

Larger dataset converges roughly 10x faster than smaller datasets for given α .

Key Observation: As n increases: $w_{\text{prior}} \rightarrow 0$ faster (exponentially in n) \Rightarrow Larger n causes faster method convergence

Part 3: Changes When Strategy 1 is Replaced by Strategy 2

Observation: The ML error curve remains invariant across strategies since it depends only on sample statistics. However, the Bayes and MAP curves show dramatically different trends: in Strategy 1, PoE generally increases with alpha, while in Strategy 2, PoE generally decreases with alpha. Despite these opposing trends, both strategies ultimately converge to similar relative performance relationships at very large alpha values.

Explanation:

The divergent behavior between strategies is rooted in the quality and appropriateness of the prior information encoded in each strategy.

The discriminant function for classification: $g(x) = x^T \Lambda^{-1} \mu - (1/2) \mu^T \Lambda^{-1} \mu + \log P(\text{class})$

PoE depends critically on class mean separation: $\Delta\mu = |\mu_{\text{FG}} - \mu_{\text{BG}}|$

For Bayesian estimation: $\Delta\mu_n = |\Sigma_0 T (\hat{\mu}_{\text{FG}} - \hat{\mu}_{\text{BG}}) + (1/n) \Sigma_n T (\mu_0, \text{FG} - \mu_0, \text{BG})|$

Strategy 1: Discriminative Prior (Good Prior Information)

Assumption: $\mu_0, \text{FG} \approx \text{true } \mu_{\text{FG}}$ and $\mu_0, \text{BG} \approx \text{true } \mu_{\text{BG}}$ Also: $\hat{\mu}_{\text{FG}} \approx \mu_0, \text{FG}$ and $\hat{\mu}_{\text{BG}} \approx \mu_0, \text{BG}$

Therefore: $(\hat{\mu}_{\text{FG}} - \hat{\mu}_{\text{BG}}) \approx (\mu_0, \text{FG} - \mu_0, \text{BG})$

Both terms in $\Delta\mu_n$ point in SAME DIRECTION: $\Delta\mu_n = |\Sigma_0 T + (1/n) \Sigma_n T| \times (\mu_0, \text{FG} - \mu_0, \text{BG})$

Analysis across α :

At Small α : Σ_0 is small, but T is large (inverse of small number) $\Delta\mu_n \approx \Delta\mu_0$ (prior provides good separation) \Rightarrow Large separation, good discrimination \Rightarrow Low PoE

At Large α : T is small (inverse of large number) $\Delta\mu_n \approx [(1/n)\sum_n T] \times (\mu_0, FG - \mu_0, BG) \rightarrow 0$ $\Delta\mu_n \rightarrow 1$
 $\hat{\mu}_{FG} - \hat{\mu}_{BG} = \Delta\mu_{ML}$ (approaches ML separation) \Rightarrow Smaller separation than at small α \Rightarrow Higher PoE (relative to small α peak)

Conclusion: $\partial\text{PoE}/\partial\alpha > 0$ (PoE INCREASES with α)

Strategy 2: Non Discriminative Prior (Poor Prior Information)

Assumption: $\mu_0, FG \approx \mu_0, BG$ (poor prior doesn't distinguish classes) But: $\hat{\mu}_{FG} \neq \hat{\mu}_{BG}$ (data provides true separation)

Therefore: $(\mu_0, FG - \mu_0, BG) \approx 0$ (prior term nearly cancelled)

The main term becomes: $\Delta\mu_n \approx |(1/n)\sum_n T(\hat{\mu}_{FG} - \hat{\mu}_{BG})|$ (only data term matters)

Analysis across α :

At Small α : w_{prior} is significant, poor prior harms discrimination $\Delta\mu_n$ is partially cancelled by bad prior \Rightarrow Reduced separation \Rightarrow High PoE

At Large α : $w_{prior} \rightarrow 0$ (prior influence vanishes) $\Delta\mu_n \approx (\hat{\mu}_{FG} - \hat{\mu}_{BG}) = \Delta\mu_{ML}$ (approaches ML separation) \Rightarrow Increased separation compared to small α \Rightarrow Lower PoE (improving toward ML)

Conclusion: $\partial\text{PoE}/\partial\alpha < 0$ (PoE DECREASES with α)

Why Opposite Trends?:

Strategy 1: Prior ALIGNs with data \Rightarrow Good at small α , approaches ML at large α Trend: PoE up (getting worse as α increases)

Strategy 2: Prior MISALIGNs with data \Rightarrow Bad at small α , approaches ML at large α Trend: PoE down (getting better as α increases)

The key insight: The trend direction depends on $\text{Sign}(\hat{\mu} - \mu_0)$

Convergence at Large α (Both Strategies):

For both strategies, as $\alpha \rightarrow \infty$: $\mu_n \rightarrow \hat{\mu}$ $T \rightarrow 0$

$$\Delta\mu_n \rightarrow |\hat{\mu}_{FG} - \hat{\mu}_{BG}| = \Delta\mu_{ML}$$

Both strategies converge to ML separation, hence:

At large α : Both strategies have similar PoE $\approx \text{PoE}_{ML}$

This explains why curves converge in plots despite opposite trends at small α .

Code

```
train_set = load('TrainingSamplesDCT_subsets_8.mat');

alpha = load('Alpha.mat');

alpha = alpha.alpha;

for strategy_idx = 1:2

    if strategy_idx == 1

        strategy = load('Prior_1.mat');

    elseif strategy_idx == 2

        strategy = load('Prior_2.mat');

    end

for dataset_idx = 1:4

    if dataset_idx == 1

        d1_BG = train_set.D1_BG;

        d1_FG = train_set.D1_FG;

    elseif dataset_idx == 2

        d1_BG = train_set.D2_BG;

        d1_FG = train_set.D2_FG;

    elseif dataset_idx == 3

        d1_BG = train_set.D3_BG;

        d1_FG = train_set.D3_FG;

    elseif dataset_idx == 4

        d1_BG = train_set.D4_BG;

        d1_FG = train_set.D4_FG;

    end
```

```

bayes_error = [];

mle_error = [];

map_error = [];

n_FG = size(d1_FG,1);

n_BG = size(d1_BG,1);

% Loop for different alpha

for alpha_idx = 1:size(alpha,2)

cov_0 = create_prior_cov(alpha(alpha_idx), strategy.W0);

% FG

d1_FG_cov = cov(d1_FG) * (n_FG-1)/n_FG;

[mu_1_FG, cov_1_FG, mu_pred_FG, cov_pred_FG] =
estimate_bayes_params(d1_FG, d1_FG_cov, cov_0, strategy.mu0_FG, n_FG);

% BG

d1_BG_cov = cov(d1_BG) * (n_BG-1)/n_BG;

[mu_1_BG, cov_1_BG, mu_pred_BG, cov_pred_BG] =
estimate_bayes_params(d1_BG, d1_BG_cov, cov_0, strategy.mu0_BG, n_BG);

% Prior

prior_FG = n_FG / (n_FG + n_BG);

prior_BG = n_BG / (n_FG + n_BG);

% Bayes-BDR

bayes_error = [bayes_error apply_classifier('bayes', mu_pred_FG,
cov_pred_FG, mu_pred_BG, cov_pred_BG, alpha_idx, dataset_idx, strategy_idx,
prior_FG, prior_BG)];

% ML-BDR

mle_error = [mle_error apply_classifier('mle',
transpose(mean(d1_FG)), d1_FG_cov, transpose(mean(d1_BG)), d1_BG_cov,
alpha_idx, dataset_idx, strategy_idx, prior_FG, prior_BG)];

```

```

% MAP-BDR

map_error = [map_error apply_classifier('map', mu_pred_FG,
d1_FG_cov, mu_pred_BG, d1_BG_cov, alpha_idx, dataset_idx, strategy_idx,
prior_FG, prior_BG) ];

end

% plot

figure('visible','off');

plot(alpha,bayes_error,alpha,mle_error,alpha,map_error);

legend('Predict','ML','MAP');

set(gca, 'XScale', 'log');

title('PoE vs Alpha');

xlabel('Alpha');

ylabel('PoE');

saveas(gcf,['Strategy_' int2str(strategy_idx) '_dataset_'
int2str(dataset_idx) '_PoEvsAlpha.png']);

close(gcf);

end

end

% Zigzag traversal

function output = zigzag(in)

h = 1;

v = 1;

vmin = 1;

hmin = 1;

vmax = size(in, 1);

hmax = size(in, 2);

```

```

i = 1;

output = zeros(1, vmax * hmax);

while ((v <= vmax) && (h <= hmax))

    if (mod(h + v, 2) == 0)                                % going up

        if (v == vmin)

            output(i) = in(v, h);                         % if we got to the first line

        if (h == hmax)

            v = v + 1;

        else

            h = h + 1;

        end

        i = i + 1;

    elseif ((h == hmax) && (v < vmax))      % if we got to the last column

        output(i) = in(v, h);

        v = v + 1;

        i = i + 1;

    elseif ((v > vmin) && (h < hmax))      % all other cases

        output(i) = in(v, h);

        v = v - 1;

        h = h + 1;

        i = i + 1;

    end

else                                % going down

    if ((v == vmax) && (h <= hmax))      % if we got to the last line

```

```

        output(i) = in(v, h);

        h = h + 1;

        i = i + 1;

elseif (h == hmin)                                % if we got to the first column

    output(i) = in(v, h);

    if (v == vmax)

        h = h + 1;

    else

        v = v + 1;

    end

    i = i + 1;

elseif ((v < vmax) && (h > hmin))      % all other cases

    output(i) = in(v, h);

    v = v + 1;

    h = h - 1;

    i = i + 1;

end

if ((v == vmax) && (h == hmax))                % bottom right element

    output(i) = in(v, h);

    break

end

end

end

% Create prior covariance

```

```

function cov_0 = create_prior_cov(alpha_val, W0)

cov_0 = zeros(64,64);

for idx = 1:64

cov_0(idx,idx) = alpha_val * W0(idx);

end

end

% Estimate Bayesian parameters

function [mu_1, cov_1, mu_pred, cov_pred] = estimate_bayes_params(data,
data_cov, cov_0, mu0, n_samples)

tmp = inv(cov_0 + (1/n_samples)*data_cov);

mu_1 = cov_0 * tmp * transpose(mean(data)) + (1/n_samples) * data_cov *
tmp * transpose(mu0);

cov_1 = cov_0 * tmp * (1/n_samples) * data_cov;

% predictive distribution (normal distribution)

mu_pred = mu_1;

cov_pred = data_cov + cov_1;

end

% Apply classifier

function total_error = apply_classifier(classifier_type, mu_FG, cov_FG,
mu_BG, cov_BG, alpha_idx, dataset_idx, strategy_idx, prior_FG, prior_BG)

% Load and process image

img = imread('cheetah.bmp');

img = im2double(img);

% Add paddle

img = [zeros(size(img,1),2) img];

img = [zeros(2, size(img,2)); img];

img = [img zeros(size(img,1),5)];

img = [img; zeros(5, size(img,2))];


```

```

[m,n] = size(img);

Blocks = ones(m-7,n-7);

% Prepare constants for discriminant function

det_cov_FG = det(cov_FG);

det_cov_BG = det(cov_BG);

ave_tmp_FG = transpose(mu_FG);

ave_tmp_BG = transpose(mu_BG);

inv_tmp_FG = inv(cov_FG);

inv_tmp_BG = inv(cov_BG);

% Pre-compute constants

const_FG = ave_tmp_FG*inv_tmp_FG*transpose(ave_tmp_FG) + log(det_cov_FG) -
2*log(prior_FG);

const_BG = ave_tmp_BG*inv_tmp_BG*transpose(ave_tmp_BG) + log(det_cov_BG) -
2*log(prior_BG);

% Classify

for i=1:m-7

    for j=1:n-7

        DCT = dct2(img(i:i+7,j:j+7));

        zigzag_order = zigzag(DCT);

        feature = zigzag_order;

        g_cheetah = compute_score(feature, inv_tmp_FG, ave_tmp_FG,
const_FG);

        g_grass = compute_score(feature, inv_tmp_BG, ave_tmp_BG,
const_BG);

```

```

    if g_cheetah >= g_grass

        Blocks(i,j) = 0;

    end

end

% Save prediction

imwrite(Blocks, [classifier_type '_prediction_alpha_' int2str(alpha_idx)
'_dataset_' int2str(dataset_idx) '_strategy_' int2str(strategy_idx) '.jpg']);

prediction = mat2gray(Blocks);

% Calculate error

total_error = calculate_error(prediction, prior_FG, prior_BG);

end

% Compute discriminant score

function score = compute_score(feature, inv_cov, ave_tmp, const)

    score = feature*inv_cov*transpose(feature) -
2*feature*inv_cov*transpose(ave_tmp) + const;

end

% Calculate classification error

function total_error_64 = calculate_error(prediction, prior_FG, prior_BG)

ground_truth = imread('cheetah_mask.bmp')/255;

x = size(ground_truth, 1);

y = size(ground_truth, 2);

count1 = 0;

count2 = 0;

count_cheetah_truth = 0;

count_grass_truth = 0;

```

```

for i=1:x

    for j=1:y

        if prediction(i,j) > ground_truth(i,j)

            count2 = count2 + 1;

            count_grass_truth = count_grass_truth + 1;

        elseif prediction(i,j) < ground_truth(i,j)

            count1 = count1 + 1;

            count_cheetah_truth = count_cheetah_truth + 1;

        elseif ground_truth(i,j) >0

            count_cheetah_truth = count_cheetah_truth + 1;

        else

            count_grass_truth = count_grass_truth + 1;

        end

    end

end

error1_64 = (count1/count_cheetah_truth) * prior_FG;
error2_64 = (count2/count_grass_truth) * prior_BG;
total_error_64 = error1_64 + error2_64;

end

```